A Deep Learning Approach to Stress Prediction Based on Sleep and Screen Time Patterns

Autor: Andrea Medina Rico **Student ID:** A01705541

Abstract. This report presents a deep learning approach for stress level prediction based on sleep patterns, screen time, and lifestyle-related features. The problem was formulated as a regression task given the continuous nature of stress. Two Multi-Layer Perceptrons (MLP) were implemented and compared through their performance. The final model achieved a Mean Absolute Error (MAE) of 0.4515 and demonstrated consistent convergence during training. These results indicate that the model is able to capture underlying patterns between behavioral features and stress levels. This work highlights the effectiveness of deep learning models for stress prediction and the importance of proper preprocessing and model design when dealing with health-related data.

I. INTRODUCTION

Nowadays, stress has become an important health concern in young adults and working populations. Despite its relevance, it is often normalized or it remains undetected until it has severe consequences.

With the use of digital devices, screen time has increased and is often extended until late hours at night. This change has been linked with negative effects in sleep patterns, deteriorating well-being. Since both sleep behavior and screen usage are measurable, they can provide signals for understanding a person's stress levels.

This report proposes a deep learning model to predict stress levels using sleep and screen time data. The goal is to contribute to early stress detection systems that can support preventive mental health interventions.

II. DATA DESCRIPTION

This dataset is called *Sleep, Screen Time and Stress Analysis*, obtained from Kaggle. It contains 15, 000 synthetic observations that aim to represent individuals with different lifestyles and habits.

Categorical features:

- **user_id:** Unique identifier per observation.
- **gender:** Male, Female or Other.
- **occupation:** Professional background of the individual.

Numerical features:

- **age:** Age of the individuals.
- **daily_screen_time_hours**: Total of hours using digital devices.
- **phone_usage_before_sleep_minutes** Minutes using digital devices before going to sleep.
- **sleep_duration_hours**: Average hours of sleep per night.
- **sleep_quality_score**: Score from 1 to 10 ranking sleep quality.
- **stress_level**: Measured stress in scale from 1 to 10
- **caffeine_intake_cups**: Caffeinated drinks during the day.
- **physical_activity_minutes**: Minutes spent on physical activity.
- **notifications_received_per_day**: Notifications from smartphone per day.
- **mental_fatigue_score**: Mental fatigue in scale from1 to 10.

III. PREPROCESSING

a. Transformation

No null observations were identified, but a column does not help the purpose of the model, as it does not provide meaningful information about the individual's behavior: the user id. Therefore, the column was removed. Additionally, instances containing missing values were excluded, as their presence could bias the model and negatively affect predictive performance.

The other two categorical features were encoded in order to become an input to the model. The *gender* feature had three values, as shown previously. Due to the low representation of the *Other* category (4%) in *gender*, it was excluded to prevent the model from bias during training. This decision is based on data limitations and is not related to gender diversity considerations. The *gender* feature turned into a binary variable, having 0 for Male and 1 for Female.

Occupation had 8 different values: Designer, Teacher, Software Engineer, Manager, Student, Freelancer, Doctor and Researcher. To prevent the model from interpreting them as ordinal values, one hot encoding was applied. Using *k - 1* dummy variables, multicollinearity between them was avoided. The value *Designer* was not included, meaning when all the dummy variables equal 0, it corresponds to *Designer*.
Except for the scores and the stress level, the numerical features were obtained in two scales: hours and minutes. To ensure a consistent scale, all of them were converted to hours. Therefore, the column names were updated to reflect the correct units.

After the data split, all the features were standardized in order to improve training stability. It helps to give all the variables the same weight. This is done after the split to prevent data leakage from the test data.

b. Data correlation

The correlation matrix shows that most of the features have weak or almost null linear relationships between themselves and the target value. This indicates that stress level is influenced by complex relationships with the features rather than just linear. Therefore, a deep learning model could help capture this pattern.

There are only three features with a relatively high linear relationship: mental fatigue score, quality sleep score and daily screen time in hours. Considering the nature of the situation, it makes sense that these have a stronger relationship with stress.

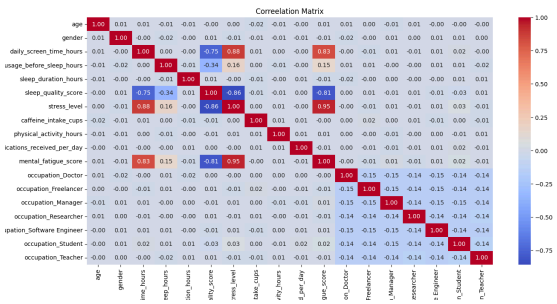

Figure 1. Correlation matrix

c. Data distribution

The target value stress level is skewed to the left with a higher concentration of observations near the maximum value. This also happens with the mental fatigue score. The rest of features have either an almost perfect balance or a normal distribution.

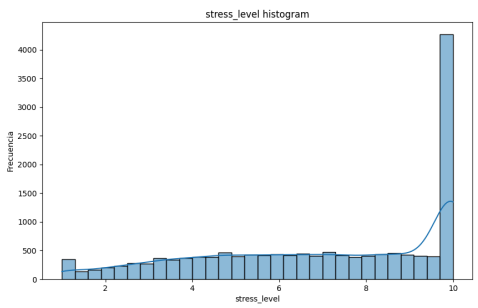

Figure 2. Stress level (target) histogram, skewed to the left.

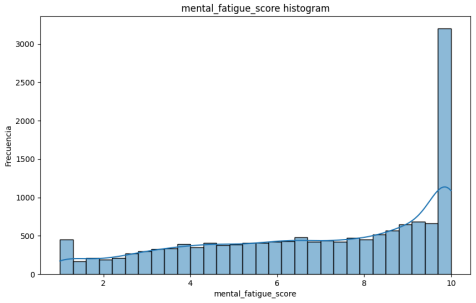

Figure 3. Mental fatigue score histogram, skewed to the left.

To improve training stability, all the features were standardized leaving them with mean equal 0 and standard deviation equal 1. This ensures that all variables contribute equally to the training. Features with a different distribution or scale don't influence more in the model. This process is done after the split, preventing data leakage from the testing instances.

d. Data split

The dataset does not have a previous separation for training and testing. Therefore, 70% of it will be dedicated to training, 15% to validation and 15% to test.

IV. MODEL

a. Justification

Prior work often formulates stress prediction as a classification problem. However, in health-related contexts, treating the target as a continuous value through a regression approach allows for precise and informative predictions [1]. In this paper, several models were compared. It was proven that Multi-Layer Perceptron (MLP) can effectively identify complex dependencies in tabular health-related datasets, improving predictive performance compared to linear approaches

The loss function was Mean Squared Error (MSE), while the metric used for interpretation was Mean Absolute Error (MAE). These were chosen because MAE measures the average magnitude of prediction errors in the units of the target value, making it understandable for a human. MSE is used as a loss function as it penalizes larger deviations more strongly due to the squared values. Using both provides a more comprehensive evaluation. Prior work in stress prediction problems has employed these evaluation metrics, supporting their suitability for this task [2].

b. Model description

A MLP will be used with the following architecture:

- Dense layer 32 neurons
- Dense layer with 16 neurons
- Single-neuron output

The choice of a relatively simple architecture is intentional, as the dataset contains structured tabular data. This does not require very deep networks to be effectively modeled, reducing the risk of overfitting.

The first two layers use ReLU as their activation function to introduce non-linearity and capture more complex relations. The learning rate optimizer is Adam, which makes the learning rate adapt between epochs. It is suitable for problems with heterogeneous feature scales, common in behavioral datasets. The performance is measured in Mean Squared Error (MSE) and Mean Absolute Error (MAE).

c. Evaluation

As a diagnosis, the model is fit, meaning it doesn't present meaningful bias or difference between training and validation. The training and validation curves show a rapid decrease during the first epochs, indicating that the model quickly learns the patterns in the data. After approximately

5–10 epochs, both curves stabilize, suggesting convergence of the model. The final MAE suggests that the stress level predictions have an average error of 0.4748.

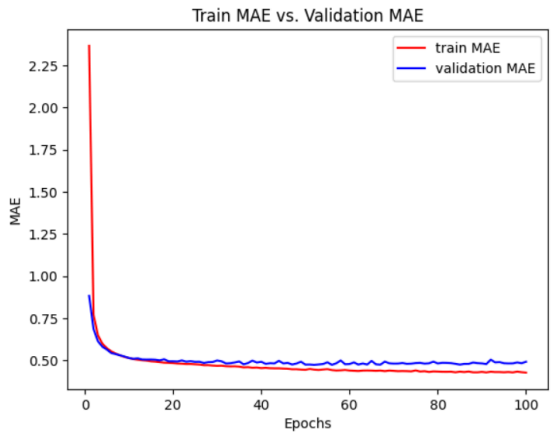

Figure 4. Train vs Validation MAE

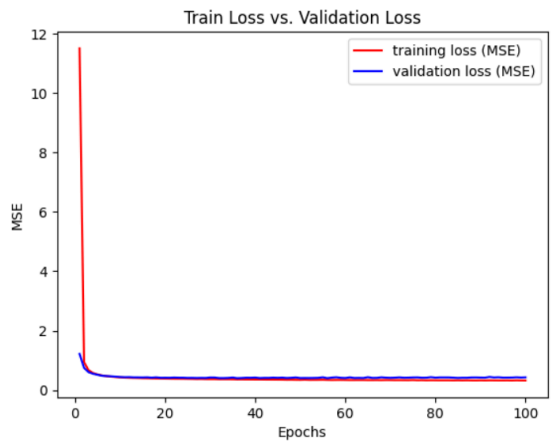

Figure 5. Train vs Validation Loss

The predicted versus actual values plot shows that most predictions are closely distributed around the ideal diagonal line where $y = x$. However, there is a considerable amount of points deviating noticeably from the reference line, with the largest errors being approximately 2 to 3 units away from the true values.

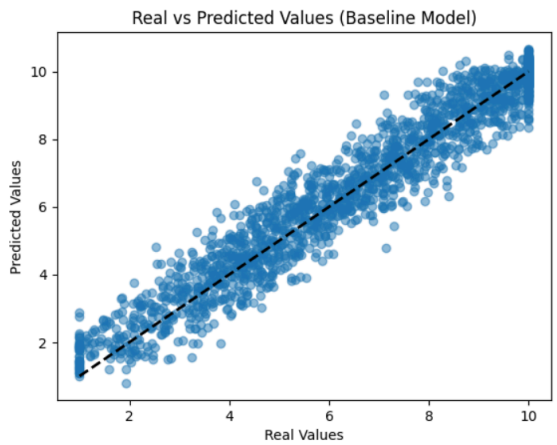

Figure 6. Real vs. Predicted Values

V. MODEL REFINEMENT

a. Justification

To improve the model's performance, the architecture will be expanded in both layers and neurons (128-64-32). Prior work on health-related tabular data has successfully employed architectures with progressively decreasing layer sizes, showing that increasing depth enables the model to better capture patterns while maintaining generalization through regularization techniques such as dropout [4]. In this work, a pattern of 128-64-32 neurons was used with a dropout of 20%.

b. Model description

The refined architecture consists of:

- Dense with 128 neurons
- Dense with 64 neurons
- Dense with 32 neurons
- Single-neuron output

After the first two dense layers, Dropout was added. As the architecture became more complex and the dataset remains as structured tabular data, Dropout helps to prevent overfitting by randomly deactivating neurons during training. This encourages the model to learn more patterns and representations.

The activation function remains ReLU and the output layer remains linear to give the continuous value as a prediction.

The Adam optimizer is used with a reduced learning rate of 0.0005 to allow more stable parameter updates during training. This adjustment helps the model converge more smoothly and leads, as shown in the posterior graphs.

The model is trained for 150 epochs, increasing the number of iterations compared to the baseline to allow sufficient learning under the new architecture. Performance is evaluated using Mean Squared Error (MSE) and Mean Absolute Error (MAE).

c. Evaluation

Even with the deeper architecture, the model does not present overfitting nor underfitting, keeping a similar performance in both train and validation. Adding more epochs helped the model to converge even more. The final MAE suggests that the stress level predictions have an average error of 0.4515.

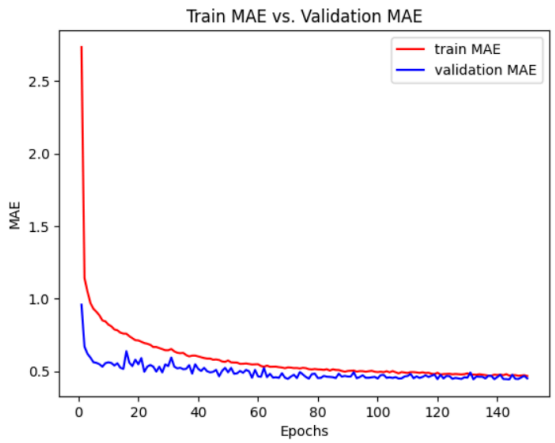

Figure 7. Train vs. Val MAE in refinement

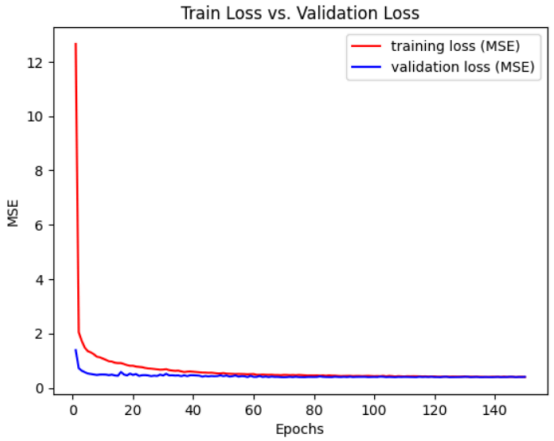

Figure 8. Train vs. Val Loss in refinement

The predicted values look closer to the line, as the color blue looks the darkest on the center. Even though some points are deviating, most of them are almost precise.

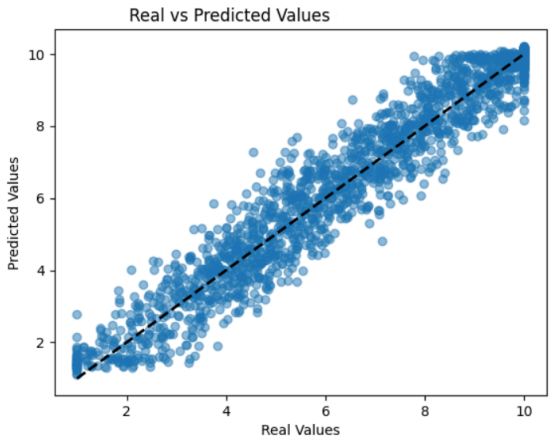

Figure 9. Real vs. Predicted values in refinement

VI. RESULTS COMPARISON

The following table compares the hyperparameters used and the metrics results for both models.

Aspect	Model 1	Model 2
Layers	3	4
Neurons	32, 16, 1	128, 64, 32, 1
Learning rate	0.001	0.0005

LR Optimizer	Adam	Adam
Dropout	No	Yes
Activation function	ReLU	ReLU
MSE	0.3970	0.4050
MAE	0.4774	0.4515

Figure 10. Table comparing hyperparameters and results between both models

The difference in MAE is 0.0259, meaning the second model is 4.32% better than the first one.

VII. CONCLUSION

This report evaluates and compares two Multilayer Perceptron architectures to predict stress level using screen time and sleep quality patterns. The baseline model demonstrated strong performance with relatively low error values, confirming that deep learning is suitable for this task.

The enhanced model (Model 2) introduced a deeper architecture along with Batch Normalization, Dropout regularization, and a reduced learning rate to improve learning stability and generalization. Although the improvement in MSE was not significant, it achieved a lower MAE, indicating more consistent predictions in terms of average deviation.

These results suggest that increasing model complexity and adding regularization techniques does not necessarily lead to significant improvements in all evaluation metrics. As the model doesn't significantly improve with more complex architectures or features, further improvements would require adding more steps to data processing.

VIII. REFERENCES

[1] Alazaidah R, Samara G, Aljaidi M, Haj Qasem M, Alsarhan A, Alshammari M. Potential of Machine Learning for Predicting Sleep Disorders: A Comprehensive Analysis of Regression and Classification Models. Diagnostics (Basel). 2023 Dec 22;14(1):27. doi: 10.3390/diagnostics14010027. PMID: 38201336; PMCID: PMC10802836.

[2] Zhang Y, Qi E (2022) Happy work: Improving enterprise human resource management by predicting workers' stress using deep learning. PLOS ONE 17(4): e0266373. https://doi.org/10.1371/journal.pone.0266373

[3] Alshayeji, Mohammad & Abed, Sa'Ed. (2025). Heart disease prediction by tabular modeling with deep learning network and interpretability. Machine Learning: Science and Technology. 6. 10.1088/2632-2153/adfd39.